
Deep Learning

Marc Lelarge, Jill-Jenn Vie, and Kevin Scaman

Class notes by Antoine Groudiev



Last modified 14th June 2024

Contents

1	Introduction and general overview	4
1.1	What is Deep Learning?	4
1.1.1	Neural networks	4
1.1.2	Timeline of Deep Learning	4
1.1.3	Recent applications and breakthroughs	4
1.1.4	Usual setup	4
1.1.5	Required skills	4
1.1.6	Building blocks of deep learning	4
1.1.7	Why deep learning now?	4
1.2	Machine Learning pipeline	4
1.2.1	Cats vs. dogs	4
1.2.2	Typical Machine Learning setup	4
1.2.3	Training objective	4
1.3	Multi-Layer Perceptron	4
1.3.1	Definition	4
1.3.2	PyTorch implementation	4
2	Automatic differentiation	4
3	Introduction to Reinforcement Learning	4
4	Optimization and loss functions	4
5	Convolutional Neural Networks	4
6	Recursive Neural Networks	4
7	Attention and Transformers	4
8	Robustness and regularity	4
9	Q-Deep Learning for Breakout	4
10	Autoencoders	4
11	Generative Adversarial Networks	4
12	Normalizing Flows	4

Abstract

This document is Antoine Groudiev's class notes while following the class *Deep Learning* at the Computer Science Department of ENS Ulm. It is freely inspired by the class notes written by Marc Lelarge, Jill-Jênn Vie, and Kevin Scaman.

1 Introduction and general overview

1.1 What is Deep Learning?

1.1.1 Neural networks

1.1.2 Timeline of Deep Learning

1.1.3 Recent applications and breakthroughs

1.1.4 Usual setup

1.1.5 Required skills

1.1.6 Building blocks of deep learning

1.1.7 Why deep learning now?

1.2 Machine Learning pipeline

1.2.1 Cats vs. dogs

1.2.2 Typical Machine Learning setup

1.2.3 Training objective

1.3 Multi-Layer Perceptron

1.3.1 Definition

1.3.2 PyTorch implementation

2 Automatic differentiation

3 Introduction to Reinforcement Learning

4 Optimization and loss functions

5 Convolutional Neural Networks

6 Recursive Neural Networks

7 Attention and Transformers

8 Robustness and regularity

9 Q-Deep Learning for Breakout

10 Autoencoders

11 Generative Adversarial Networks

12 Normalizing Flows